

# Democratic Spaces 2022

Updated forecasts and a refinement in how opening and closing events are measured

Andreas Beger (Basil Analytics)\*

26 May 2022

## Contents

<b>Summary</b>	<b>1</b>
<b>ERT-lite: new method for identifying opening and closing events</b>	<b>2</b>
<b>Overview of the new 2022-2023 forecasts</b>	<b>3</b>
<b>Accuracy for past forecasts</b>	<b>4</b>
Shifting ground . . . . .	6
Accuracy for v9 and v10 forecasts, partial assessment for v11 . . . . .	7
<b>Summary of other technical changes and additional resources</b>	<b>8</b>
Dashboard changes . . . . .	8
Data and model changes . . . . .	9
ERT-lite Assessment . . . . .	9
<b>Additional Resources</b>	<b>11</b>
<b>References</b>	<b>11</b>

## Summary

- The way that opening and closing events in the six spaces has changed. The new method (“ERT-lite”) is able to identify more gradual changes in the spaces that over the span of two or more years result in significant shifts. The old method was limited to identifying large single-year shifts—which are still included in the new method as well. This allows DemSpaces to, for example, pick up the more gradual closings in some of the spaces for Poland and Hungary over the 2010s.
- The new forecasts overall indicate higher risks of continued closing events, reflecting the ongoing global worsening of democratic governance. While on balance the trend leans towards

---

\*Email: [andy@basilanalytics.com](mailto:andy@basilanalytics.com)

negative, there are also still significant potentials for opening movements in other countries and spaces.

- The accuracy of the v10 forecasts, which can now be fully assessed, is in line with the previous v9 accuracy. With AUC-ROC values of 0.7 the forecasts could be more accurate—0.8 or slightly higher is common in many political instability applications—but are informative additions over a naive baseline estimate. There is some indication that the stability of the most recent V-Dem version has increased, which would alleviate the shifting-ground problem induced by changes in historical data as each V-Dem version continues to refine the data.


(Note: the forecasts are indexed with the V-Dem data version used to produce them. For the latest, v12 forecasts, I use “v12” to denote forecasts using the old outcome coding method, and “v12.1” for the new ERT-lite outcome coding method. Both sets of forecasts can be found in the project repo [demspaces/archive](#) folder, but only the “v12.1” forecasts are shown in the dashboard.)

## ERT-lite: new method for identifying opening and closing events

Opening and closing events in the democratic spaces have so far been based on large *single-year* changes. However, sometimes there are cases where the space in a country changes gradually but significantly over a span of multiple years, without ever experiencing a large single-year change that our coding approach to date would have identified. Figure 1 shows the governing space in Poland. Since 2014 it has been experiencing gradual but steady declines. With the old outcome coding method, none of those years were identified as closing events.

**Figure 1:** Poland governing space: there is a large, gradual decrease since 2014, but no closing event is coded in the old method.

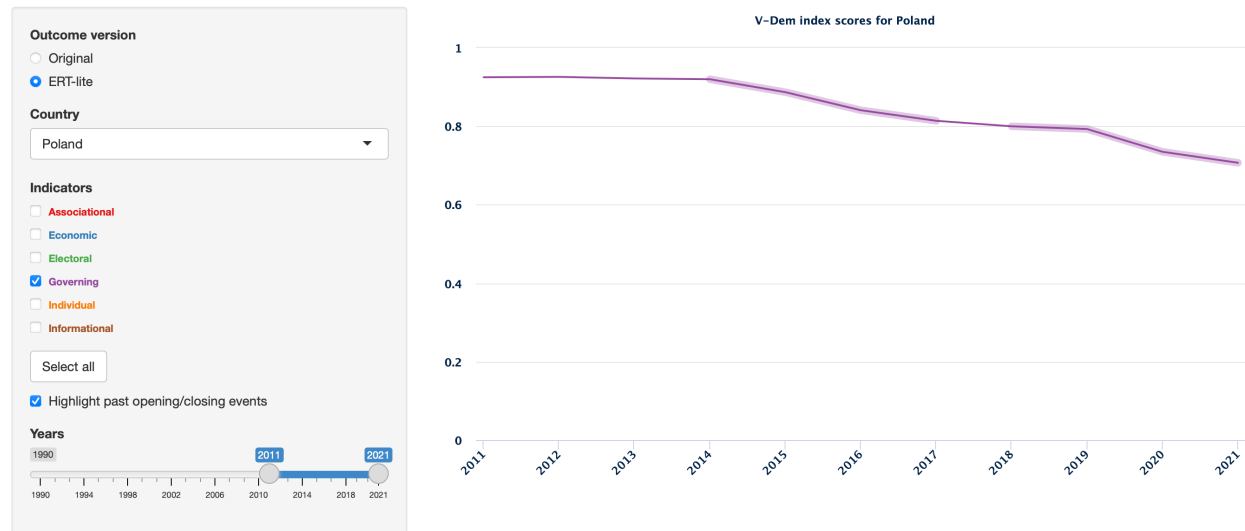
### Democratic Spaces Outcome Explorer



As part of the spring 2022 update, I investigated alternative outcome coding methods that would better capture gradual transitions like this. The end result, for the Poland governing case, is that most of the years during the gradual transition since 2014 are now coded as such, as shown in Figure 2.

**Figure 2:** Poland governing space with the new outcome coding method: most of the years during the gradual closing since 2014 are identified as closing events.

### Democratic Spaces Outcome Explorer



The new method for coding opening and closing events was inspired by V-Dem’s Episodes of Regime Transition project (Maerz et al. 2021), which unifies autocratization and democratization in a single conceptual framework (the eponymous “episodes of regime transformations”) and describes an algorithm for identifying relevant cases in the V-Dem electoral democracy index. Their project and algorithm are much broader in scope than what is required for democratic spaces, but it spurred the improvement in the opening and closing cases coding method that I showed above.

Namely, to date opening and closing events are coded if the change in a space from the previous year exceeds a space-specific threshold value. Each space has a different threshold because some spaces are in general more fluctuating than others. To this single-year rule I added another two-year rule: if the cumulative 2-year change in a space exceeds the space threshold, *and* each year’s change from the previous year exceeds 1/10th of the threshold (in the same direction), then both years are coded as an being opening or closing events, depending on the direction of change.

This is a minimally intrusive extension—we go from only considering a single-year change to considering a two-year change—but manages to capture gradual transitions better than the existing approach. Table 1 shows the overall change in the number of opening and closing events for 2021. In all but one instance—Electoral openings—the number of events increases.

More details of this change are discussed in Section “ERT-lite Assessment” below.

## Overview of the new 2022-2023 forecasts

Because the new, ERT-style outcome coding method identifies more opening and closing events than the old method, we should expect that the forecasts generally are higher. But there may at the same time also be systemic changes in the global risk of opening and closing events, e.g. increasing instability.

**Table 1:** Comparison of old- and new-style outcomes

Space	Closing		Same		Opening	
	v12	v12.1	v12	v12.1	v12	v12.1
Associational	20	28	135	123	14	18
Economic	21	26	128	120	20	23
Electoral	7	9	156	154	6	6
Governing	18	28	143	129	8	12
Individual	14	26	146	127	9	16
Informational	20	30	143	125	6	14

To untangle this, Figure 3 summarizes last year’s “v11” forecasts, the directly comparable new “v12” forecasts using the old-style outcome coding method, and the new ERT-lite “v12.1” forecasts. The only difference between the “v12” and “v12.1” forecasts is that the latter uses the new ERT-lite outcome coding method. We can see that the forecasts in general are indeed higher—from an average of 14.0% to 19.4% and 13.8% to 18.5% for closing and opening, respectively. By comparing the “v12” and “v11” forecasts, where the only difference is that the former is using updated data, we can assess whether there are global trends in the democratic spaces. There are no dramatic changes. The average opening risk is unchanged at 13.8% and the average closing risk slightly increases from 13.4% to 14.0%.


**Figure 3:** Distribution of last year’s and this year’s forecasts, pooled over spaces.

Table 2 shows the 5 highest forecasts for opening and closing changes for each of the 6 spaces (using the “v12.1” forecasts).

## Accuracy for past forecasts

There are now four sets of forecasts, corresponding to the four V-Dem data versions since version 9 in 2019. These are shown in Table 3.

**Table 2:** Highest risk for opening and closing events in each space.

Closing		Opening	
Country	Pr.	Country	Pr.
<b>Associational</b>			
India	0.74	Bangladesh	0.69
Philippines	0.73	Somalia	0.65
Hungary	0.69	Pakistan	0.58
Somalia	0.64	Nicaragua	0.57
Thailand	0.63	South Korea	0.52
<b>Economic</b>			
Brazil	0.68	Saudi Arabia	0.88
Indonesia	0.60	China	0.76
Colombia	0.60	Gabon	0.64
India	0.56	Angola	0.60
Lithuania	0.55	Nepal	0.60
<b>Electoral</b>			
Poland	0.52	Mali	0.69
Guinea	0.46	Qatar	0.59
Afghanistan	0.40	Guinea	0.50
Chad	0.39	Sudan	0.38
El Salvador	0.36	Tunisia	0.37
<b>Governing</b>			
Poland	0.68	Turkey	0.80
United Arab Emirates	0.64	Angola	0.62
Indonesia	0.62	Uzbekistan	0.55
Somalia	0.62	Ecuador	0.54
Canada	0.54	Paraguay	0.52
<b>Individual</b>			
Brazil	0.76	Ukraine	0.55
Bangladesh	0.74	Congo	0.55
Qatar	0.74	Sierra Leone	0.55
Poland	0.67	Uzbekistan	0.54
India	0.66	Venezuela	0.54
<b>Informational</b>			
Brazil	0.81	Benin	0.62
India	0.75	Bhutan	0.61
Philippines	0.72	Pakistan	0.59
Hungary	0.71	Somalia	0.57
Indonesia	0.69	Ukraine	0.55

**Table 3:** Overview of DemSpaces forecasts to date

Version	Last year	Forecasts	Can be scored
v9	2018	2019 – 2020	X
v10	2019	2020 – 2021	X
v11	2020	2021 – 2022	partial
v12.1	2021	2022 – 2023	

**Table 4:** Agreement rates between opening and closing cases in different V-Dem data versions.

Comparison	Agreement Rate
v9-v10	0.66
v10-v11	0.68
v11-v12	0.88
v9-v11	0.63
v10-v12	0.66

The first two sets of forecasts can be fully assessed as enough new data has accumulated to measure outcomes for their two-year forecast periods. We can partially score last year's (v11) forecasts with 2021 outcomes.

## Shifting ground

Forecast accuracy is influenced by changes in the democratic space indicator variables between V-Dem versions.<sup>1</sup> In practice this has the effect of decreasing forecast accuracy.

Fortunately, changes in the outcomes between V-Dem data versions are decreasing.

Table 4 shows year to year (first 3 rows), and year to year + 2 (this is how the forecasts are assessed) agreement rates for opening and closing outcome cases. This is for 1970 to 2018, when the v9 data ended. The "Agreement Rate" column shows agreement rates for cases that are a positive in either dataset version, i.e. the number of joint opening and closing events that *both* dataset versions coded as opening or closing.

The first three rows are year-to-year changes. For the first two years of DemSpaces, it was around 0.65. The outcomes with this year's v12 data however match 0.88 with the v11 data, much higher. That should mean that the accuracy of the current forecasts will be higher than it has been for the v9 and v10 forecasts.

Table 5 shows what all cases look like between v11 and v12<sup>2</sup>. This is with the full data from 1970 to 2021. More than 80% of opening and closing cases are coded as such in both dataset versions. When there are coding differences, it is always a opening or closing case that is considered as no change ("Same") in the other dataset. Never do the two datasets code changes in opposing directions.

<sup>1</sup>This shifting-ground phenomenon is described in more detail in the 2021 update report, "[Democratic Spaces Dashboard: 2021 Update and Accuracy Assessments](#)".

<sup>2</sup>Note that I am using the old-style outcome coding. That is the one relevant for scoring the older forecasts.

**Table 5:** Cross-tabulation of v11 and v12 outcomes

v11	v12		
	Closing	Same	Opening
Closing	524	68	0
Same	55	9937	40
Opening	0	57	461

**Table 6:** Accuracy of the v9 (2019) forecasts, scored with v11 data.

Space	Cases	High Risk	AUC-ROC	AUC-PR	Pos. Rate
<b>Closing</b>					
Associational	29	12	0.65	0.25	0.17
Economic	31	13	0.72	0.30	0.18
Electoral	10	6	0.72	0.12	0.06
Governing	21	7	0.62	0.15	0.12
Individual	34	14	0.70	0.36	0.20
Informational	30	11	0.67	0.26	0.18
<b>Opening</b>					
Associational	16	8	0.75	0.22	0.09
Economic	42	17	0.67	0.42	0.25
Electoral	4	1	0.64	0.03	0.02
Governing	26	11	0.72	0.24	0.15
Individual	19	12	0.82	0.29	0.11
Informational	19	10	0.77	0.35	0.11
<b>Average</b>	23	10	0.71	0.25	0.14

Let's go back to the last two rows in Table 4. Since the forecasts are two years ahead, it takes two V-Dem data updates to fully observe the outcomes. Thus v9 forecasts are scored with v11 data, not v10 data. The second two rows in the table show the agreement rates for opening and closing cases with a two year version difference. It has also been increasing, although there still is a lot of shifting ground that the forecasts have to measure against.

Generally, the higher this agreement is, the higher the accuracy of the forecasts will be. The test forecasts have for example AUC-ROC values greater than 0.8, so with more agreement the actual performance should converge to some value like that.

The good news is that the agreement rate has been becoming higher, indicating less changes in the historical data between updated versions of V-Dem.

### Accuracy for v9 and v10 forecasts, partial assessment for v11

Table 6 shows the accuracy of the v9 (2019 - 2020) forecasts, scored with v11 data. This is the same table as in the report from the 2021 update last year.

The accuracy of the v10 (2020 - 2021) forecasts, scored with v12 data, is shown in Table 7. It is

**Table 7:** Accuracy of the v10 (2020) forecasts, scores with v12 data

Space	Cases	High Risk	AUC-ROC	AUC-PR	Pos. Rate
<b>Closing</b>					
Associational	33	13	0.69	0.28	0.20
Economic	41	14	0.62	0.31	0.24
Electoral	10	5	0.72	0.11	0.06
Governing	29	11	0.71	0.29	0.17
Individual	30	11	0.70	0.29	0.18
Informational	29	11	0.68	0.24	0.17
<b>Opening</b>					
Associational	24	12	0.75	0.28	0.14
Economic	39	16	0.66	0.30	0.23
Electoral	7	5	0.83	0.12	0.04
Governing	20	8	0.74	0.19	0.12
Individual	17	6	0.66	0.13	0.10
Informational	13	4	0.64	0.19	0.08
<b>Average</b>	24	10	0.70	0.23	0.14

overall comparable to the v9 forecast performance.

Table 8 shows the **partial** accuracy of the v11 (2021 - 2022) forecasts from last year, with outcomes for 2021 but missing for 2022. These numbers will go up when data for this year becomes available. Right now they are lower than the accuracy for the v9 and 10 forecasts.

Overall, there is no change in the assessment of accuracy from last year's report: the DemSpaces forecasts are limited but informative indicators of opening and closing potential. If the V-Dem data become more stable, as they appear to be with the latest update, then there is reason to expect that the overall performance will converge to a level with AUC-ROC values of around 0.8, which is consistent with the performance in a variety of conflict forecasting applications.

## Summary of other technical changes and additional resources

This section reviews some other changes that occurred during this year's update.

### Dashboard changes

- Added functionality to highlight past opening and closing events in the time series plot on the bottom right of the dashboard. ([Issue #12](#))
- Minor dashboard design changes to make it easier to select/de-select all spaces on the time-series plot ([Issue #13](#)); better text layout for wide screens ([Issue #20](#)); various updates to the About tab ([Issue #21](#)); and overall changes to the design of the dashboard text, header formatting, etc. ([Issue #24](#)).



**Table 8:** Partial accuracy of the v11 (2021) forecasts, scores with v12 data

Space	Cases	High Risk	AUC-ROC	AUC-PR	Pos. Rate
<b>Closing</b>					
Associational	20	8	0.65	0.18	0.12
Economic	21	6	0.57	0.13	0.12
Electoral	7	4	0.70	0.08	0.04
Governing	18	8	0.69	0.23	0.11
Individual	14	6	0.61	0.12	0.08
Informational	20	7	0.69	0.17	0.12
<b>Opening</b>					
Associational	14	8	0.76	0.16	0.08
Economic	20	8	0.60	0.15	0.12
Electoral	6	4	0.82	0.09	0.04
Governing	8	5	0.80	0.10	0.05
Individual	9	5	0.71	0.11	0.05
Informational	6	3	0.73	0.06	0.04
<b>Average</b>	14	6	0.69	0.13	0.08

## Data and model changes

- Eliminated impossible forecasts, e.g. when a country has such as high value in a space already that it cannot increase beyond the relevant space threshold to have an opening event. This is fixed by hard-coded post-processing of the raw model forecasts. ([Issue #15](#))
- Removed “\_squared” variable transforms and instead added a moving 10-year standard deviation as indicator of general recent instability (“\_sd10”); this overall slightly improved model accuracy. ([Issue #18](#))
- Tuning experiments to review and find better hyperparameter values for the random forest models. All hyperparameters are now fixed, which speeds up model estimation. The tuning experiments are described in the “[RF tuning experiments](#)” note. The number of trees in the models is set at an intentionally higher value to reduce the random variation of point forecasts, see the “[RF Stability](#)” note.

## ERT-lite Assessment

The new outcome coding algorithm is:

1. If a single-year change exceeds the relevant space threshold  $cp_s$ , code as opening, and if it exceeds  $-cp_s$  code as closing event. (This is what we already had.)
2. If the cumulative change in a 2-year window exceeds the relevant space threshold  $cp_s$  (the same threshold we already have), and both years exceed  $0.1 \times cp_s$ , code both years as opening, and similarly for the closing direction.

In addition to the Poland-Governing example that is shown above, I used Hungary Associational and Informational as test cases during the algorithm development. (These can be seen at the GitHub repo in [issue # 16](#).) It turned out that the minimal extension above satisfactorily addressed the three test cases I used, and thus I did not explore further modifications.

**Table 9:** Impact of modified outcome coding on opening/closing rates. All values are percentages. 1970 – 2021, N = 8,187.

Space	Opening			Closing		
	Orig	Mod	Increase	Orig	Mod	Increase
Informational	4.8	8.9	88	3.1	6.4	107
Governing	4.4	6.9	58	2.9	5.1	75
Economic	4.7	5.9	25	5.6	6.6	19
Electoral	3.1	4.8	55	2.2	3.2	46
Individual	4.7	8.7	87	3.0	6.0	97
Associational	5.8	8.8	53	3.8	6.3	64

One concern with the original ERT algorithm is that it results in a potentially very long time period before it is clear whether a country-year was in an episode or not. Long feedback cycles are problematic for forecast model development and assessment. The simple extension adds just 1 year to the time required to fully score one of the DemSpaces 2-year forecasts. For example, the current 2022-2023 forecasts will be fully realized with the spring 2025 V-Dem update instead of the spring 2024 V-Dem update.

Another potential concern was on forecast model accuracy. The new coding increases positive opening and closing cases quite a bit—although both are overall still uncommon events. Table 9 shows the overall impact of the modified outcome coding (“v12.1” data, the original is “v12” of the data) on the number of opening and closing events in the full data from 1970 – 2021. All values are shown as rates (percentages). There are 8,187 total country-year cases in the data, so 1% represents roughly 82 cases.

The number of country-years coded as being in an opening or closing episode increases by 66% on average, from a rate of roughly 4% to 6.5% of all country-years. Opening and closing events are thus overall still rare situations.

To assess the potential risk of reduced forecast model performance, I used cross-validation to obtain out-of-sample predictions for all cases we have in the data, from 1970 - 2021. To do this I split the data from 1970 to 2019<sup>3</sup> into 5 or 10 groups<sup>4</sup>, and then for each group, used the remaining 4 or 9 groups to estimate two models, one for the original outcome and the other for the modified outcomes. These two models then predict outcomes in the held-back group, creating out-of-sample predictions. By doing this for all groups we can obtain out-of-sample predictions for the entire data. I repeated the cross-validation several times with different group splits, giving an overall sample of N = 30 out-of-sample fit statistic values.

Table 2 shows the overall performance, pooled across spaces and directions. The first row shows the positive rate for the indicator that the models are forecasting—whether an opening or closing episode occurred during the subsequent 2 years.<sup>5</sup>

The new outcome coding does not seem to have a negative impact on performance. AUC-ROC

<sup>3</sup>The end date was determined by when the modified outcomes are fully resolved.

<sup>4</sup>The splitting was done by year, i.e. I randomly assigned years to folds.

<sup>5</sup>The rates are different from those in Table 9 because we are here additionally summarizing over 2 years. E.g. if only one year in the next 2 was part of an opening episode, it overall gets coded as a 1, thus essentially amplifying the positive rate.

**Table 10:** Cross-validation out-of-sample performance using both the original and modified outcome coding.

Measure	Modified (v12.1)	Original (v12)
Pos-rate	0.10	0.07
PR-AUC	0.29	0.23
ROC-AUC	0.81	0.81

is essentially the same. AUC-PR increases slightly, but part of that is because the positive rate increases, i.e. the prediction problem becomes slightly easier.

## Additional Resources

The Democratic Spaces Dashboard can be found at: <https://www.v-dem.net/demspace>

The data, code, and additional resources can be found in the GitHub repo at: <https://github.com/vdeminstitute/demspaces>. This includes:

- An overview of recent changes in “[NEWS.md](#)”.
- Additional notes related to several of the technical changes mentioned above, see the “[2022-update](#)” folder.
- A note investigating the high risk of Associational closing in France, “[What-if France](#)”.

## References

Maerz, Seraphine M., Amanda B. Edgell, Matthew C. Wilson, Sebastian Hellmeier, and Staffan I. Lindberg. 2021. “A Framework for Understanding Regime Transformation: Introducing the ERT Dataset.”